



GUIDED TOUR

Dashboard Walkthrough

A tour of the product in the order that lands every differentiator: the command center, one call end-to-end, the quality/measurement story, and the run-from-scratch guarantee.

Setup


```
docker compose up -d          # db + api + web
docker compose ps             # confirm all 3 are healthy
# → dashboard http://localhost:3000 · API http://localhost:8000
```

If the database is empty on a fresh machine, `./scripts/restore.sh` brings the whole analysed dataset back in seconds — no re-transcription.

The 30-second opening

"Every bank records its support calls 'for quality' — and nobody listens to them. A manager who owns 300 calls a day can review maybe five. Companion for CallRadar listens to all 1,441. It turns raw recordings into the five things a manager needs — who called, what they wanted, whether their mood turned, whether it was really resolved — and **the exact second on the call that proves every one of those claims.**"

Beat 1 — The Command Center



1,441 calls at a glance.

"Nine need a manager today. But look — **195 calls that sounded resolved but weren't**. That's the number no manual review would ever find. Evidence coverage is 100% — every judgment is backed by a verified quote. And across the whole team, agents verified customer identity on only 7% of sensitive calls — a compliance finding worth the entire product."

Beat 2 — One call, end to end

CallRadar
Conversation Intelligence

Overview

Needs Attention

QA & Compliance

Customers

Agents

Trends

Quality

← Back

Jennifer Rodriguez

Agent David · 0:48 · 3c4891908c3d4498 · engine c:claude;c:claude-opus-4-8

Recording

AI summary ≤ 40 words

Jennifer called to replace a lost credit card. The agent confirmed which card but never verified her identity or explicitly confirmed a replacement was ordered, then abruptly ended the call without a timeframe.

lost credit card · card replacement · identity verification gap

Mood timeline

Neutral

Neutral

Judgments every claim cites the moment that justifies it — click to hear it

INTENT — WHAT THEY WANTED

Replace a lost credit card · card lost stolen

00:17 · customer ✓ verified

"I lost my credit card. Can you send me a new one?"

MOOD — STEADY

Neutral

00:14 · customer ✓ verified

"Hi, my name is Jennifer Rodriguez."

RESOLUTION

Unclear

The agent confirmed the card type and closed the call, but never verified identity, confirmed a replacement was actually ordered, or provided a delivery timeframe. No explicit confirmation the card was reissued.

00:42 · agent ✓ verified

"Great. Thanks very much. Bye."

NEEDS-ATTENTION SCORE

65/100

Lost card report closed with no identity verification

No explicit confirmation a replacement card was ordered

Agent ended call abruptly without delivery timeframe

Possible security/process gap on a card-loss request

00:42 · agent ✓ verified

"Great. Thanks very much. Bye."

Evidence verification: 5/5 cited quotes matched the transcript at the stated timestamp.

QA & Compliance how well the call was handled — every check cites the moment

22/100

⚠️ RESOLUTION RISK — sounded resolved but wasn't

• a sensitive action (card / password / transfer) was taken without verifying the customer's identity

• the call closed politely but the requested action was never confirmed as completed

AGENT OPENED PROFESSIONALLY

PASS

00:01 · agent ✓ verified

"Hello. Hello, this is Harper Valley National Bank. My name is David. How can I help you today?"

VERIFIED CUSTOMER IDENTITY BEFORE A SENSITIVE ACTION

FAIL

A sensitive action (card/password/transfer) happened with no identity check. This is a compliance and fraud risk.

CONFIRMED THE ACTION WAS COMPLETED

FAIL

Agent never confirmed the request was actually done — the customer may leave unsure.

ACKNOWLEDGED / SHOWED EMPATHY

FAIL

Acknowledge the customer's situation before solving.

DID NOT HAVE TO REPEAT THE SAME QUESTION

FAIL

00:20 · agent ✓ verified

"Yeah, thank you for letting me know. Let me see if I can help you replace your card. Which card would you like to replace? Your debit or credit card? Credit card? Perfect. And is there anything else I"

CLOSED THE CALL PROPERLY

PASS

00:20 · agent ✓ verified

"Yeah, thank you for letting me know. Let me see if I can help you replace your card. Which card would you like to replace? Your debit or credit card? Credit card? Perfect. And is there anything else I"

Agent asked the same thing more than once — listen actively / avoid rework.

Transcript 7 turns · agent = left channel, customer = right channel

Recording, transcript, evidence, mood timeline, QA.

1. **Recording + transcript** — raw audio in, turn-by-turn out. Agent left channel, customer right.

2. **The evidence** — click any judgment's timestamp: it jumps the audio to that exact second and highlights the line.
 3. **The summary** — 40 words, and specific: it caught that the agent moved money but never verified identity.
 4. **QA & Compliance** — this call is marked *Resolved* with a positive mood. It *sounds* fine. But it's flagged: money transferred with no identity verification — a fraud and compliance risk. Every check cites its moment.
-

Beat 3 — "How do we know it's right?"

"Most teams show a working demo. We show it's *measured*: **faithfulness 99.9%** (we verify each quote actually appears at the cited time — this is how we defend against 'wrong evidence scores negative'), **diarization 100% separation** (channel-split makes attribution correct by construction), and **coverage 100%**."

Beat 4 — Run it yourself

"It all runs from `docker compose up`, then one pipeline command turns the recordings into everything you've seen. **No API keys required** — there's a fully offline engine. Plug in Claude or Azure OpenAI and the top calls get LLM-grade analysis. Transcription never re-runs on a request; it's all precomputed in Postgres."

Questions you'll get, and the answers

QUESTION	ANSWER
How is "who said what" so accurate?	Stereo channels — agent left, customer right. We split and transcribe separately, so attribution is correct by construction, not guessed.
What if the model makes up evidence?	We verify every quote against the transcript at its timestamp. Unverified ones show in amber. Measured faithfulness is 99.9%.
What's your accuracy?	We deliberately don't quote a brittle exact-match number — the research shows it punishes good-but-reworded answers. We measure faithfulness, diarization, and coverage.
Does it need an API key?	No — it runs fully offline. Add Claude or Azure OpenAI to upgrade the calls that matter.
The "resolution risk" — is that real?	Yes. It flags calls that closed politely but were never confirmed done, or moved money without identity verification. Every flag cites the moment.

Companion for CallRadar — Walkthrough. See also Architecture, Formulas, the complete documentation.